

# midterm

## 01 Intro

- NLP's goal
- Common Applications
- Interpretation of garden-path sentences

Regex	Match (single characters)	Example Patterns Matched
/[A-Z]/	not an upper case letter	"Oyfn pripetchik"
/[Ss]/	neither 'S' nor 's'	"I have no exquisite reason for't"
/[.]/	not a period	"our resident Djinn"
/[e']/	either 'e' or 'e'	"look up _ now"
/a^b/	the pattern 'a^b'	"look up a^b now"

/a\*/ means "any string of zero or more as"

a+ means "one or more occurrences of the immediately preceding character or regular expression"

^ matches the start of a line. The pattern /^The/ matches the word The only at the start of a line. \$ end of line

\b matches a word boundary, and \B matches a non-boundary

The pattern /cat|dog/ matches either the string cat or the string dog.() means and a\.{24}z/ will match a followed by 24 dots followed by z (but not a followed by 23 or 25 dots followed by a z)

Regex	Expansion	Match	First Matches
\d	[0-9]	any digit	Party_of_5
\D	[^0-9]	any non-digit	Blue_moon
\w	[a-zA-Z0-9_]	any alphanumeric/underscore	Daiyu
\W	[^\w]	a non-alphanumeric	!!!!
\s	[\r\t\n\f]	whitespace (space, tab)	in_Concord
\S	[^\s]	Non-whitespace	in_Concord

{n} exactly n occurrences of the previous char or expression

{n,m} from n to m occurrences of the previous char or expression

{n,} at least n occurrences of the previous char or expression

{,m} up to m occurrences of the previous char or expression \n a newline \t a tab

In that case we use a non-capturing group, which is specified by putting the special commands ?: after the open parenthesis, in the form (?: pattern ).

The operator (?! pattern) only returns true if a pattern does not match, but again is zero-width and doesn't advance the cursor.

Herdan's Law (Herdan, 1960) or Heaps' Law (Heaps, 1978)  $|V| = kN^\beta$

Two events, A and B, are independent iff:  $P(A, B) = P(A)P(B)$

word tokens - an individual word instance. (a list) word types - distinct words. (a set)

V - "vocabulary" |V| - vocabulary size (number of types) N - number of tokens

Corpus - a natural language dataset

- Tokenization ■ Parts of speech ■ word tokenizer ■ word piece tokenization (conceptually)

## BPE

– more data-driven; no predefined words or rules 如何合并的

– allow for subwords (e.g. "unlikeliest" -> "un", "like", "liest") – better for unseen words or capturing semantics of parts of words.

```
function BYTE-PAIR ENCODING(strings C, number of merges k) returns vocab V
  V ← all unique characters in C      # initial set of tokens is characters
  for i = 1 to k do                  # merge tokens k times
    t_L, t_R ← Most frequent pair of adjacent tokens in C
    t_NEW ← t_L + t_R                # make new token by concatenating
    V ← V + t_NEW                    # update the vocabulary
    Replace each occurrence of t_L, t_R in C with t_NEW    # and update the corpus
  return V
```

## 2 3 Maximum Entropy Classifier Supervised Machine Learning

Build a “model” that can estimate  $P(Y=1|X=?)$   $Y=1$  if target is a verb

使用逻辑回归的函数作为decision boundary  $L(\beta|X, Y) = \prod_{i=1}^n P(Y_i = 1|x_i)^{y_i} (1 - P(Y_i = 1|x_i))^{1-y_i}$

best fit: whatever maximizes the likelihood function  $J(\beta) = -\frac{1}{N} \sum_{i=1}^N (y_i \log p(x_i) + (1 - y_i) \log(1 - p(x_i)))$

## Logistic Regression on a single feature (x)

$Y_i \in \{0, 1\}$ ;  $X$  is a single value and can be anything numeric.

$$P(Y_i = 1 | X_i = x) = \frac{e^{\beta_0 + \beta_1 x_i}}{1 + e^{\beta_0 + \beta_1 x_i}} = \frac{1}{1 + e^{-(\beta_0 + \sum_{j=1}^m \beta_j x_{ij})}}$$

$Y_i \in \{0, 1\}$ ;  $X$  is a **single value** and can be anything numeric.

$$P(Y_i = 1 | X_i = x) = \frac{1}{1 + e^{-(\beta_0 + \sum_{j=1}^m \beta_j x_{ij})}}$$

Vector notation

$\beta$  and  $x_i$  are vectors of size  $m$

first feature is intercept:  
 $x_{*,0} = [1, 1, \dots, 1]_N$

$$= \frac{1}{1 + e^{-(x_i \beta)}}$$

We're still learning a linear separating hyperplane, but fitting it to a logit outcome.应用逻辑函数将线性预测结果转换成对数几率，得到分类概率，不改变决策边界

$\beta \approx \text{weight} \approx \text{coefficient} \approx \text{parameters} \approx \theta$  Logistic Regression  $\approx$  Maximum Entropy Classifier loss function  $\approx$  cost function

```
z = torch.matmul(x, beta) #将输入向量x与参数向量beta进行矩阵乘法运算，得到一个线性预测结果z。这里假设x是一个输入特征向量，beta是一个参数向量，二者的维度兼容。
yhat = nn.functional.relu(z) #对线性预测结果z进行ReLU（修正线性单元）激活函数的操作，得到预测输出yhat。ReLU函数将所有负值变为零，并保持非负值不变。
loss = nn.MSELoss(yhat, torch.Tensor(y)) #计算预测输出yhat与目标值y之间的均方误差（Mean Squared Error）。这里假设y是目标值，通过将其转换为`torch.Tensor`类型，与预测输出yhat进行比较计算损失值。

sgd = torch.optim.SGD(model.parameters(), lr=learning_rate)
loss_func = torch.mean(-torch.sum(y*torch.log(y_pred)))
#training loop:
for i in range(epochs):
    model.train()
    sgd.zero_grad()
    #forward pass:
    ypred = model(X)
    loss = loss_func(ypred, y)
    #backward: /(applies gradient descent)
    loss.backward()
    sgd.step()
    if i % 20 == 0:
        print(" epoch: %d, loss: %.5f" % (i, loss.item()))
```

`torch.nn.BCELoss()` binary cross entropy =  $\text{loss}(x, y) = -[y * \log(x) + (1 - y) * \log(1 - x)]$

## Machine Learning: How to setup data

o Likelihood function to Loss function row of features; e.g.

→ number of capital letters

→ whether "I" was mentioned or not

→ k features indicating whether k words were mentioned or not

Feature extraction: one-hot, multi-hot representations

### Multi-hot Encoding

• Each word gets an index in the vector • 1 if present; 0 if not

Multiple One-hot encodings for one observation

(1) word before; (2) word after (3) percent capitals

L1 Regularization - "The Lasso" Zeros out features by adding values that keep from perfectly fitting the data.

set betas that maximize penalized L  $L(\beta_0, \beta_1, \dots, \beta_k | X, Y) = \prod_{i=1}^n p(x_i)^{y_i} (1 - p(x_i))^{1-y_i} - \frac{1}{C} \sum_{j=1}^m |\beta_j|$

C is the hyperparameter, set betas that maximize penalized L

L2 Regularization - "Ridge" Shrinks features by adding values that keep from perfectly fitting the data.

$L(\beta_0, \beta_1, \dots, \beta_k | X, Y) = \prod_{i=1}^n p(x_i)^{y_i} (1 - p(x_i))^{1-y_i} - \frac{1}{C} \sum_{j=1}^m \beta_j^2$  Sometimes writing as  $\|\beta_j\|_2^2$

After doing them, we can avoid overfitting.

Overfitting is when a model is more accurate than the held-out data. --True

discriminative learning and generating learning

2. Which of the following are true about discriminative learning as compared to generative: (8 points)

- ☐ A generative model can be used to generate features supposing it is of some class c.
- ☐ A discriminative model is often used to directly create new data.
- ☐ A model which assigns classes c to a document d by seeking to estimate  $P(c|d)$  is an example of a generative modeling approach.
- ☐ Logistic regression is an example of a discriminative modeling approach.
- ☐ A model which assigns classes c to a document d by seeking to compute a likelihood term  $P(d|c)$  and a prior  $P(c)$  is an example of a discriminative modeling approach.

answer: A D

3. Recall the L2 regularized likelihood function below.

$$L(\beta_0, \beta_1, \dots, \beta_k | X, Y) = \prod_{i=1}^n p(x_i)^{y_i} (1 - p(x_i))^{1-y_i} - \frac{1}{C} \sum_{j=1}^m \beta_j^2$$

When using L2 regularization for logistic regression with penalty parameter c, which of the following statements are true?

- ☐ Each beta is guaranteed to shrink
- ☐ As c increases, betas on average increase
- ☐ As c increases, betas on average decrease
- ☐ Logistic regression will find the value of c which maximizes the likelihood function over your data

Answer: BD

4. Suppose you have the following vocabulary of words in alphabetical order for a one-hot encoding scheme.

```
[
  "I",      # 0
  "NLP",    # 1
  "dislike", # 2
  "doing",  # 3
  "like",   # 4
  "skipping" # 5
]
```

```
[
  "I",      # 0
  "NLP",    # 1
  "dislike", # 2
  "doing",  # 3
  "like",   # 4
  "skipping" # 5
]
```

You are going to produce a feature vector that concatenates the one-hot encoding of the word before and after a target word.

[1-hot of word before target] + [1-hot of word after target]

Given the sentence "I like doing NLP" select the indices below that should contain a 1 (be "hot") in the specified input feature vector if the target word is "doing".

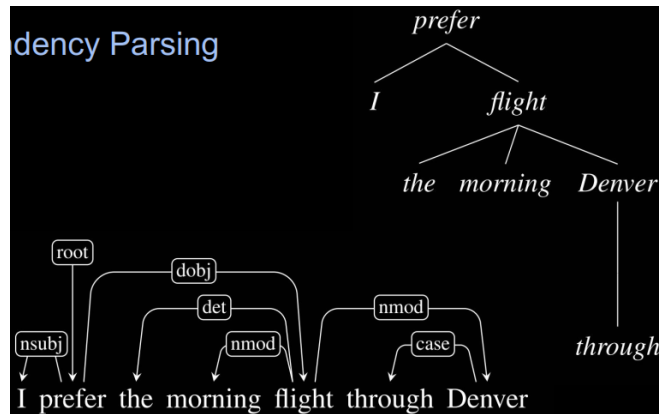
0 # I || 1 # NLP || 2 # dislike || 3 # doing || 4 # like || 5 # skipping || 6 # I || 7 # NLP || 8 # dislike || 9 # doing || 10 # like || 11 # skipping

## 4 Dependency Parsing

- Relations (core universal dependency relations) ○ head and dependent
- Transition-based dependency parsing ○ Projectivity ○ Idea of semantic roles and verbal predicates

- 关系树和依存关系 和算法

## Dependency Parsing

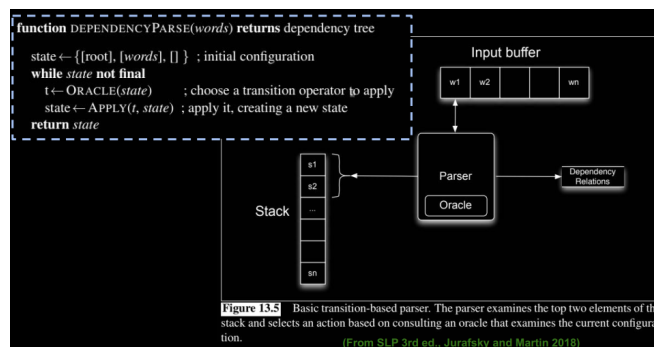


Transition-based (Shift-Reduce algorithm)

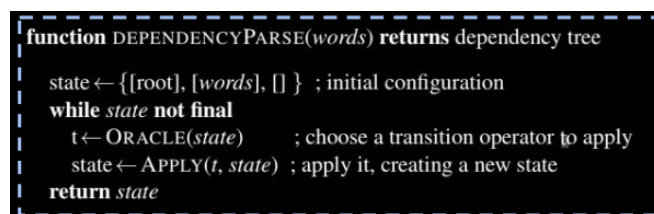
Restrictions:

1. Single designated ROOT with no incoming arcs
2. Every vertex only has one head (parent, governor); i.e. only one incoming arc
3. unique path from ROOT to every vertex

## Transition-based Dependency Parsing



- S: stack, initialized with "ROOT"
- B: input buffer, initialized with tokens ( $w_1, w_2, \dots$ ) of sentence
- A: set of dependency arcs, initialized empty
- T: Actions, given  $w_i$  (next token in stack)




From Syntax to Semantics

- We've already seen words have many meanings. ◦ Context is key
- Verbs can be seen as functions (predicates) that take arguments. ◦ Syntactic arguments fulfill semantic roles
- Words have implicit syntactic relationships with each other in given sentences.
  - Dependency Parsing: each word has one head
  - Easily constructed through 3 actions of shift-reduce parsing.

Takeaway: There is an interplay between word meaning and sentence structure!

## Graph-based Approaches

Search through all possible trees and pick best.



For each word, pick the most likely head. Then check if still a fully-connected tree, and adjust.  
Complex and slow but leads to state of the art. Now done with neural models.

## Semantic Roles

Roles are restricted to nouns, but signalled through the verb and other parts of speech.

## 5 Lexical and Vector Semantics

- terminology (lemmas, homonymy, etc...) ◦ different types of word sense disambiguation
- Lesk algorithm ◦ distributional hypothesis ◦ concept of vector semantics
- word2vec - skip-gram model ◦ topic modeling – LDA

### Question 3 (8 points) ✓ Saved

Recall the Lesk algorithm, suppose you had following senses and glosses for "cone".

1. a solid body which narrows to a point.
2. a light-sensitive cell in the eye.
3. the fruit of pine trees.

Consider the sentence "the pine cone fell". If using Jaccard Similarity, what would be the maximum overlap s

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

Round your answer to tenths places.

Your Answer:

Answer

2/7 gloss the dictionary definition of a word.

part of speech: a category to which a word is assigned in accordance with its syntactic functions.

distributional hypothesis: a word's meaning is defined by all the different contexts it appears in (i.e. how it is "distributed" in natural language).

word sense: a discrete representation of one aspect of the meaning of a word.

lemma: the canonical form, dictionary form, or citation form of a set of word forms.

### • Lexical Ambiguity (why word sense disambiguation)

Word Sense Disambiguation (WSD) `f (sent_tokens, (target_index, lemma, POS)) -> word_sense`

Distributional hypothesis -- A word's meaning is defined by all the different contexts it appears in (i.e. how it is "distributed" in natural language).

Firth, 1957: "You shall know a word by the company it keeps"

## Approaches to WSD

1. Bag of words for context --E.g. multi-hot for any word in a defined "context".
2. Surrounding window with positions --E.g. one-hot per position relative to word).
3. Lesk algorithm --E.g. compare context to sense definitions.

**function SIMPLIFIED LESK(word, sentence) returns best sense of word**

```

best-sense ← most frequent sense for word
max-overlap ← 0
context ← set of words in sentence
for each sense in senses of word do
  signature ← set of words in the gloss and examples of sense
  overlap ← COMPUTEOverlap(signature, context)
  if overlap > max-overlap then
    max-overlap ← overlap
    best-sense ← sense
end
return(best-sense)

```

**Figure 19.10** The Simplified Lesk algorithm. The COMPUTEOverlap function returns the number of words in common between two sets, ignoring function words or other words on a stop list. The original Lesk algorithm defines the context in a more complex way.

有问题在于

- 单词释义更长的单词有可能和他的使用场景有更多overlap
  - 有可能对应了意思，但是并没有overlap（相同的）单词
4. Selectors -- other target words that appear with same context  
E.g. counts for any selector. 找到同义字 Sets of selectors tend to vary extensively by word sense
  5. Contextual Embeddings --E.g. real valued vectors that "encode" the context (TBD).

## • Word Vectors

word2vec Principal: Predict missing word. Similar to classification where  $y$  = context and  $x$  = word.

2 Versions of Context:

1. Continuous bag of words (CBOW): Predict word from context
2. Skip-Grams (SG): predict context words from target
  1. Treat the target word and a neighboring context word as positive examples.
  2. Randomly sample other words in the lexicon to get negative samples
  3. Use logistic regression to train a classifier to distinguish those two cases --assume  $\text{dim} * |\text{vocab}|$  weights for each of  $c$  and  $t$ , initialized to random values (e.g.  $\text{dim} = 50$  or  $\text{dim} = 300$ )
  4. Use the weights as the embeddings

$x = (\text{hit, beam}), y = 1$   
 $x = (\text{the, beam}), y = 1$   
 $x = (\text{behind, beam}), y = 1$   
 $\dots$   
 $x = (\text{happy, beam}), y = 0$   
 $x = (\text{think, beam}), y = 0$   
 $\dots$

**single context:**  
 $P(y=1 | c, t) = \frac{1}{1 + e^{-t \cdot c}}$   
**all contexts**  
 $P(y=1 | c, t) = \prod_{i=1}^n \frac{1}{1 + e^{-t \cdot c_i}}$

*The nail hit the beam behind the wall.*

$\underbrace{\quad\quad}_c \underbrace{\quad}_t \underbrace{\quad\quad\quad}_c \underbrace{\quad}_c$   
 $c1 \quad c2 \quad \quad c3 \quad c4$

单一上下文 简化计算 固定窗口

Intuition:  $t \cdot c$  is a measure of **similarity**: But, it is not a probability! To make it one, apply logistic activation:  $\sigma(z) = 1/(1 + e^{-z})$   
 $c$  context 里面出现的词 有可能是single的有可能是all,  $t$  target(如何相乘)

Maximizes similarity of  $(c, t)$  in positive data ( $y = 1$ ) Minimizes similarity of  $(c, t)$  in negative data ( $y = 0$ )

Logistic Regression Likelihood:  $L(\beta_0, \beta_1, \dots, \beta_k | X, Y) = \prod_{i=1}^n p(x_i)^{y_i} (1 - p(x_i))^{1-y_i}$

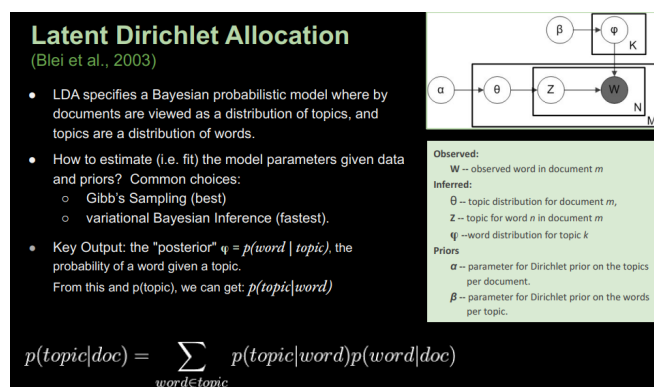
Log Likelihood:  $\ell(\beta) = \sum_{i=1}^N y_i \log p(x_i) + (1 - y_i) \log (1 - p(x_i))$

Log Loss:  $J(\beta) = -\frac{1}{N} \sum_{i=1}^N y_i \log p(x_i) + (1 - y_i) \log (1 - p(x_i))$

Cross-Entropy Cost:  $J = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^{|V|} y_{i,j} \log p(x_{i,j})$  (a "multiclass" log loss)

In vector algebra form:  $-\text{mean}(\text{sum}(y * \log(y\_pred)))$

## • Topic Modeling



Common applications:

- Open vocabulary content analysis: Describing the latent semantic categories of words or phrases present across a set of documents
- Embeddings for predictive task: for all topics, use  $p(\text{topic} | \text{document})$  as score. Feed to predictive model (e.g. classifier).
- PCA-Based Embeddings -- try to represent with only p' dimensions --also known as "Latent Semantic Analysis
- SVD-Based Embeddings -- Dimensionality Reduction - PCA

## 6 Introduction to Language Modeling

- 2 task versions and their **equivalence**
- applications
- chain rule, markov assumption
- unigram, bigram LMs -- assigning a probability to sequences of words.

Version 1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

:probability of a sequence of words --even the Web isn't large enough to enable good estimates of most phrases

Version 2: Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

:probability of a next word given history

$$P(B|A) = P(B, A) / P(A) \quad A, B = P(A)P(B|A) \Leftrightarrow P(A)P(B|A) = P(B, A) \quad P(A, B, C) = P(A)P(B|A)P(C|A, B)$$

**The Chain Rule**:  $P(X_1, X_2, \dots, X_n) = P(X_1)P(X_2|X_1)P(X_3|X_1, X_2) \dots P(X_n|X_1, \dots, X_{n-1})$

$$P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | X_1, X_2, \dots, X_{i-1})$$

**Markov Assumption**:  $P(X_n | X_1, \dots, X_{n-1}) \approx P(X_n | X_{n-k}, \dots, X_{n-1})$  where  $k < n$

Unigram Model:  $k = 0$ ; Bigram Model:  $k=1$ ;

**Perplexity** Inverse of probability (i.e. minimize perplexity = maximize likelihood) and (weighted) average branching metric for scoring how well learned model works on test. (an intrinsic evaluation)

**MLE** An intuitive way to **estimate probabilities** is called maximum likelihood estimation or MLE.

Practical Consideration:

- Use log probability for assessing perplexity to keep numbers reasonable and save computation. (uses addition rather than multiplication)
- Use Out-of-vocabulary (OOV)

Choose minimum frequency or total vocabulary size and mark as **<OOV>**

- Sentence start and end: **<s> this is a sentence </s>**

Advantage: models word probability at beginning or end.