

# Introduction to Language Modeling

H. Andrew Schwartz

CSE538 - Spring 2024



Stony Brook University  
Human Language Analysis Beings

# Language Modeling

-- assigning a probability to sequences of words.

# Language Modeling

-- assigning a probability to sequences of words.

Version 1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$   
:probability of a sequence of words

# Language Modeling

-- assigning a probability to sequences of words.

Version 1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$   
:probability of a sequence of words

Version 2: Compute  $P(w_5 | w_1, w_2, w_3, w_4)$   
 $= P(w_n | w_1, w_2, \dots, w_{n-1})$   
:probability of a next word given history

# Language Modeling

Version 1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

:probability of a sequence of words

$$P(\text{He ate the cake with the fork}) = ?$$

Version 2: Compute  $P(w_5 | w_1, w_2, w_3, w_4)$

$$= P(w_n | w_1, w_2, \dots, w_{n-1})$$

:probability of a next word given history

$$P(\text{fork} | \text{He ate the cake with the}) = ?$$

# Language Modeling

Version 1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

:probability of a sequence of words


$P(\text{He ate the cake with the fork}) = ?$

Version 2: Compute  $P(w_5 | w_1, w_2, w_3, w_4)$

$= P(w_5 | w_1, w_2, w_3, w_4)$

:probability of a next word  $g$

$P(\text{fork} | \text{He ate the cake})$




# Language Modeling

## Applications:

- Auto-complete: What word is next?
- Machine Translation: Which translation is most likely?
- Spell Correction: Which word is most likely given error?
- Speech Recognition: What did they just say?  
“eyes aw of an”

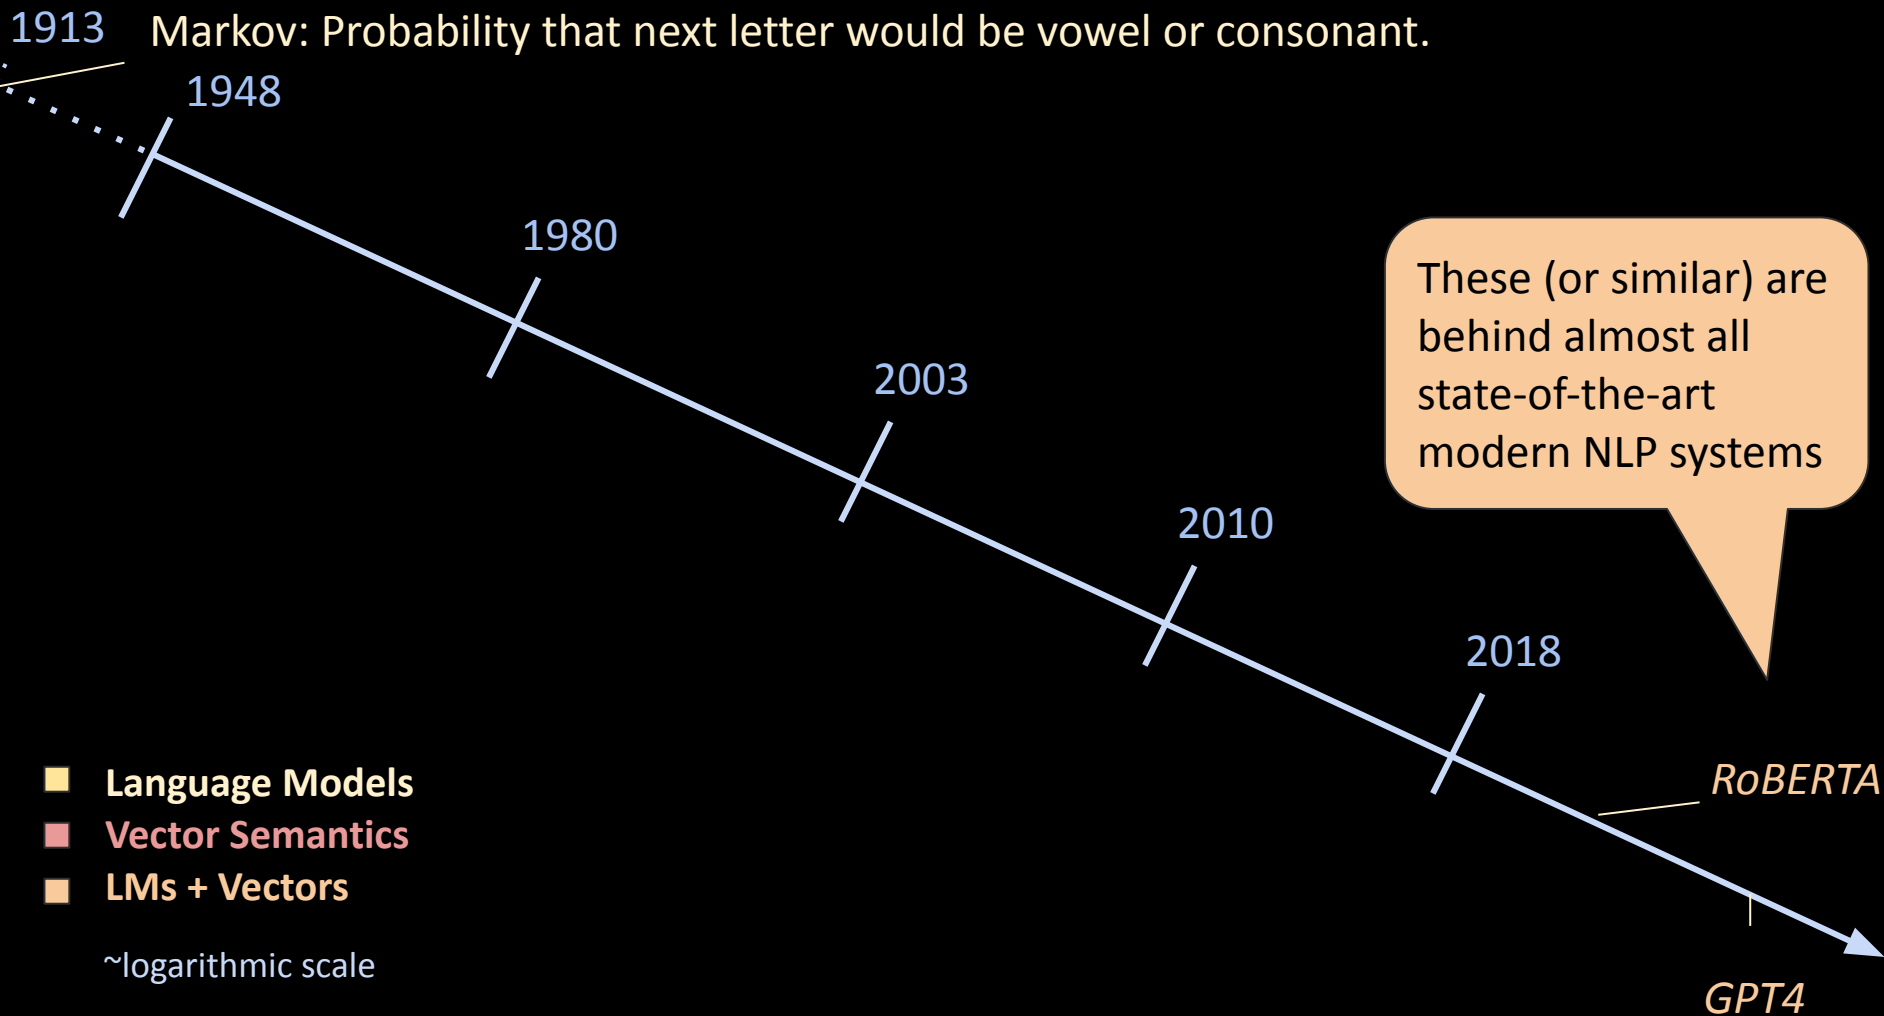
(example from Jurafsky, 2017; *did you say "giraffe ski 2,017"?*)

# Timeline: *Language Modeling* and *Vector Semantics*







# Timeline: *Language Modeling* and *Vector Semantics*



# Timeline: *Language Modeling* and *Vector Semantics*



# Timeline: *Language Modeling* and *Vector Semantics*



# Timeline: *Language Modeling* and *Vector Semantics*

1913 Markov: Probability that next letter would be vowel or consonant.

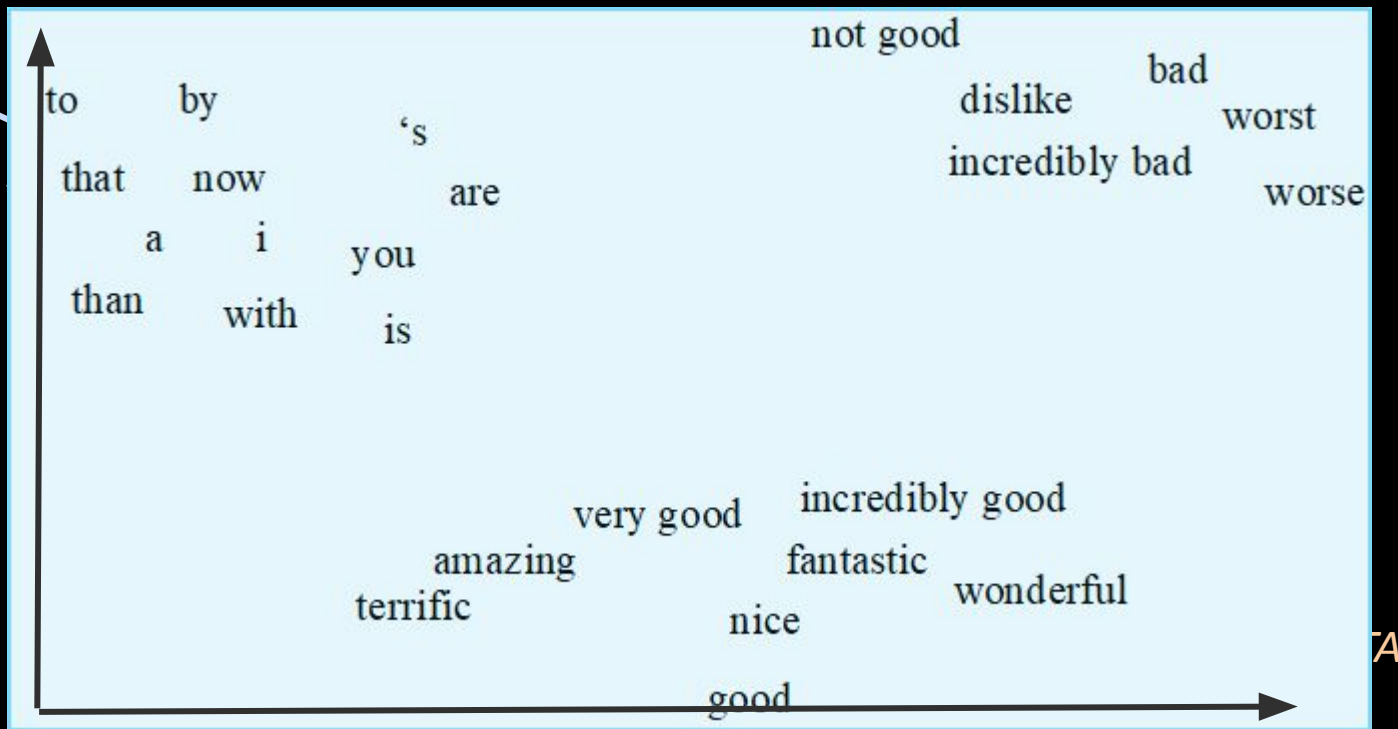
1948

Shannon: *A Mathematical Theory of Communication* (first digital language model)

Osgood: *The Measurement of Meaning*

- Language Models
- Vector Semantics
- LMs + Vectors

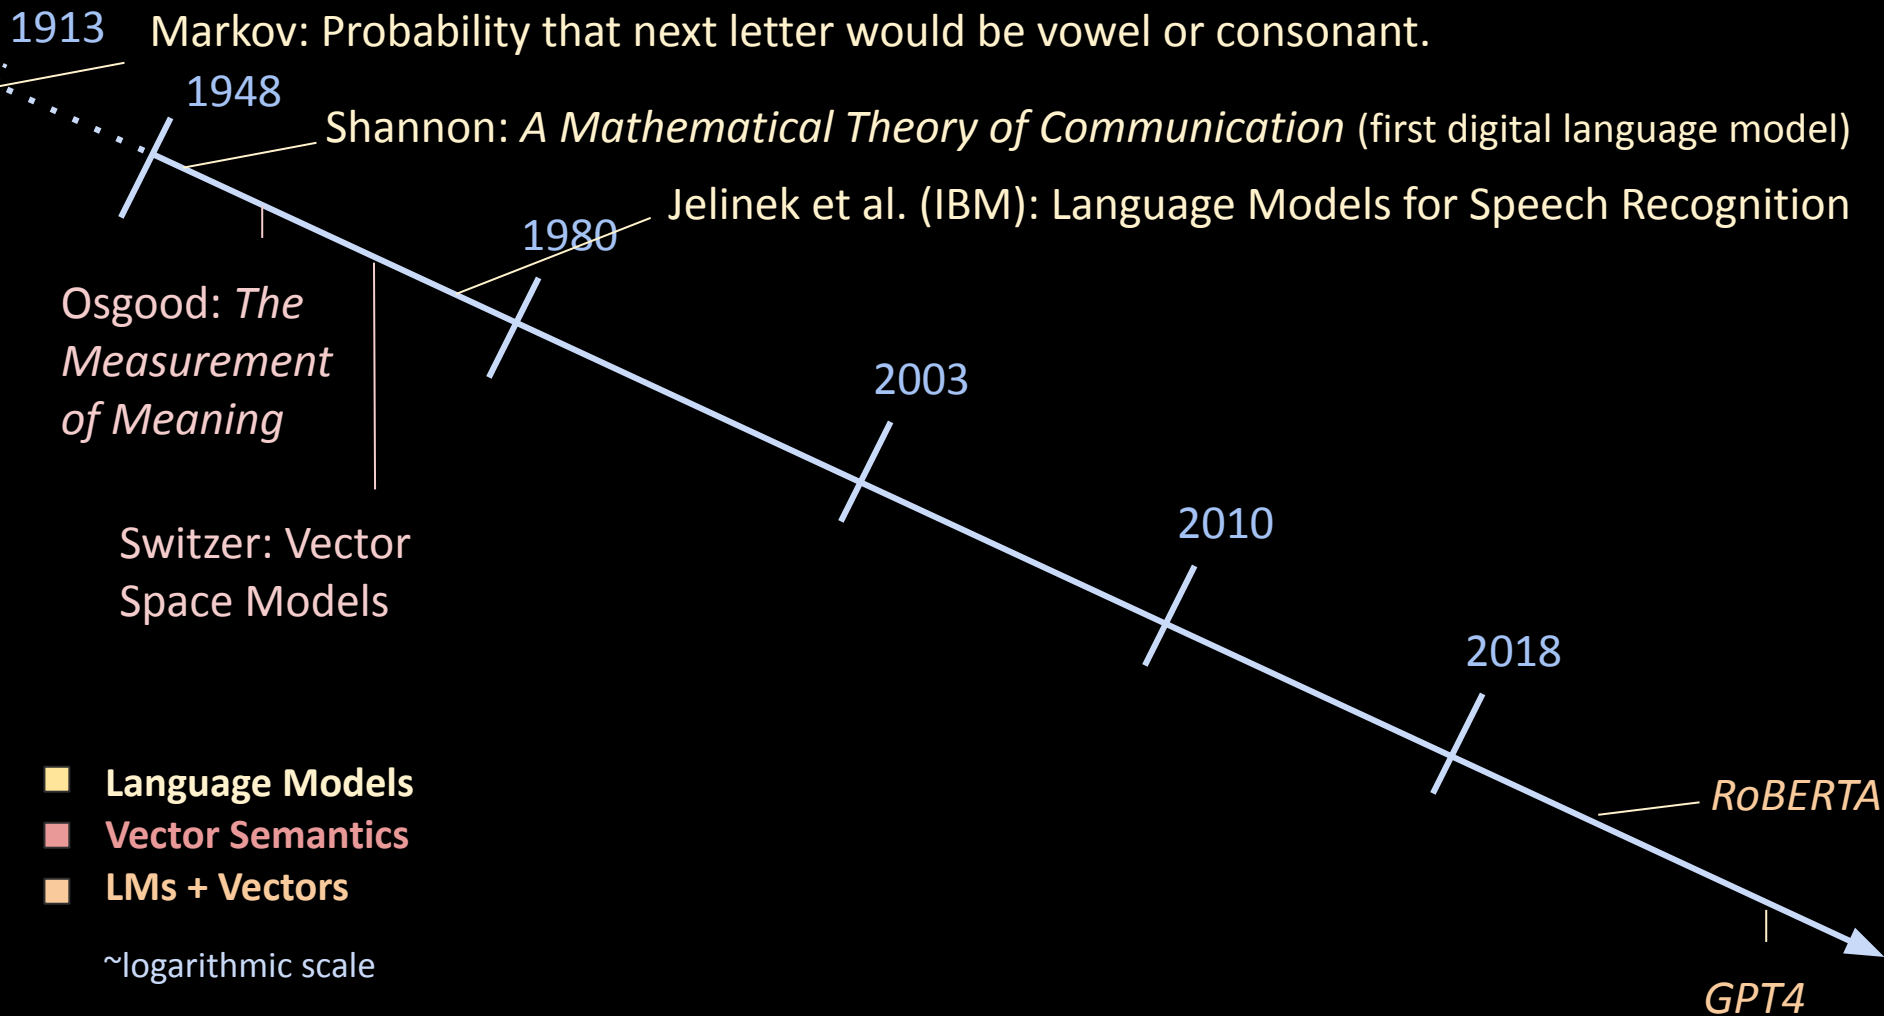
~logarithmic scale



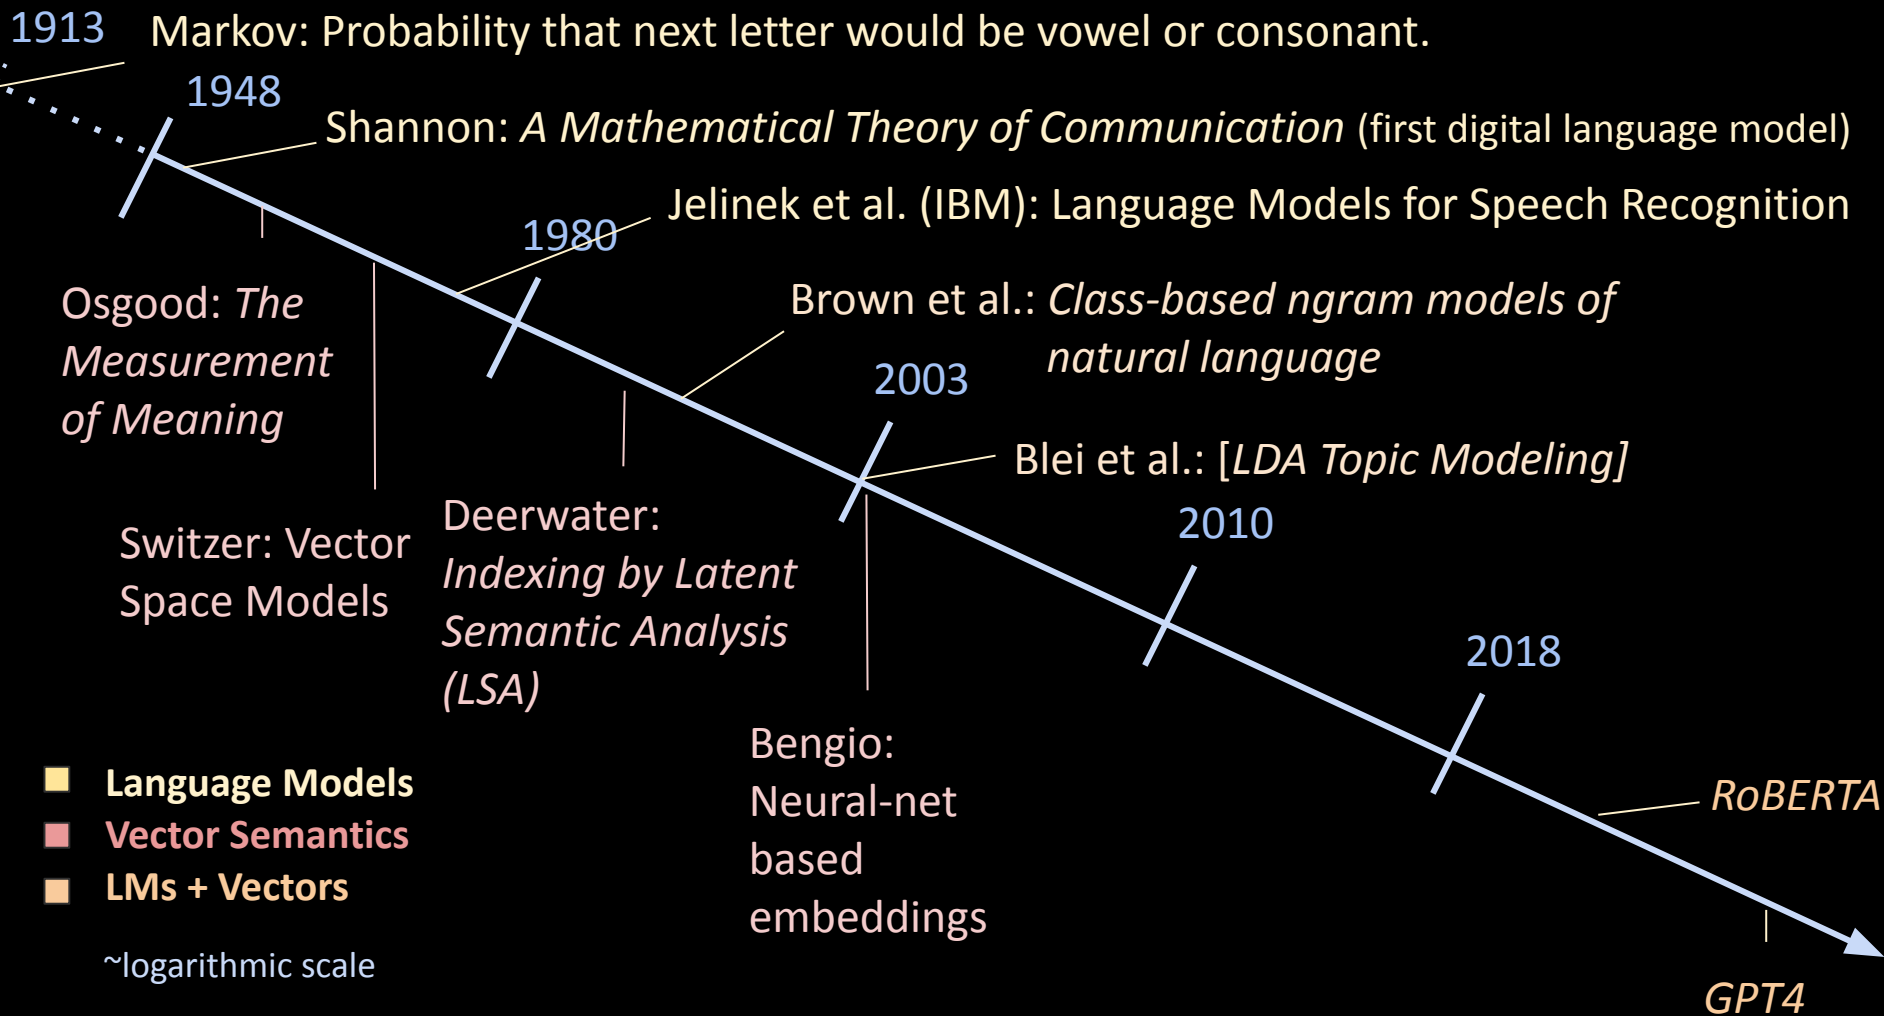
(Li et al., 2015; Jurafsky et al., 2019)

GPT3.5


# Timeline: *Language Modeling* and *Vector Semantics*




# Timeline: *Language Modeling* and *Vector Semantics*



# Timeline: *Language Modeling* and *Vector Semantics*



# Timeline: *Language Modeling* and *Vector Semantics*





# Language Modeling

Version 1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

:probability of a sequence of words

Version 2: Compute  $P(w_5 | w_1, w_2, w_3, w_4)$

$$= P(w_n | w_1, w_2, \dots, w_{n-1})$$

:probability of a next word given history

# Simple Solution

Version 1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

:probability of a sequence of words

$P(\text{He ate the cake with the fork}) =$

$$\frac{\text{count}(\text{He ate the cake with the fork})}{\text{count}(* * * * * * *)}$$

# Simple Solution: The Maximum Likelihood Estimate

Version 1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

:probability of a sequence of words

$P(\text{He ate the cake with the fork}) =$

total number of  
observed *7grams*

$$\frac{\text{count}(\text{He ate the cake with the fork})}{\text{count}(* * * * * *)}$$

# Simple Solution: The Maximum Likelihood Estimate

V1:

$$P(\text{He ate the cake with the fork}) =$$

$$\frac{\text{count}(\text{He ate the cake with the fork})}{\text{count}(* * * * *)}$$

V2:

$$P(\text{fork} \mid \text{He ate the cake with the}) =$$

$$\frac{\text{count}(\text{He ate the cake with the fork})}{\text{count}(\text{He ate the cake with the } *)}$$

# Simple Solution: The Maximum Likelihood Estimate

**Problem:** even the Web isn't large enough to enable good estimates of most phrases.

V1:

$$P(\text{He ate the cake with the fork}) =$$

$$\frac{\text{count}(\text{He ate the cake with the fork})}{\text{count}(* * * * * *)}$$

V2:

$$P(\text{fork} \mid \text{He ate the cake with the}) =$$

$$\frac{\text{count}(\text{He ate the cake with the fork})}{\text{count}(\text{He ate the cake with the } *)}$$

# Simple Solution: The Maximum Likelihood Estimate

**Problem:** even the Web isn't large enough to enable good estimates of most phrases.

V1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

V2: Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

**A solution:** Estimate from shorter sequences.

# Simple Solution: The Maximum Likelihood Estimate

**Problem:** even the Web isn't large enough to enable good estimates of most phrases.

V1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

V2: Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

**A solution:** Estimate from shorter sequences.

*Observation: V1 and V2 are equivalent!*

# Language Modeling: How to Estimate?

V1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

V2: Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

*Observation: V1 and V2 are equivalent!*



# Language Modeling: How to Estimate?

V1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

V2: Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

*Observation: V1 and V2 are equivalent!*

$$P(B|A) = P(B, A) / P(A) \Leftrightarrow P(A)P(B|A) = P(B, A)$$

# Language Modeling: How to Estimate?

V1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

V2: Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

*Observation: V1 and V2 are equivalent!*

$$P(B|A) = P(B, A) / P(A) \Leftrightarrow P(A)P(B|A) = P(B, A) = P(A, B)$$

$$P(A, B) = P(A)P(B|A)$$

# Language Modeling: How to Estimate?

V1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

V2: Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

*Observation: V1 and V2 are equivalent!*

$$P(A,B) = P(A)P(B|A)$$

# Language Modeling: How to Estimate?

V1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

V2: Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

*Observation: V1 and V2 are equivalent!*

$$P(A, B) = P(A)P(B|A)$$

$$P(A, B, C) = P(A)P(B|A)P(C|A, B)$$

# Language Modeling: How to Estimate?

V1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

V2: Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

*Observation: V1 and V2 are equivalent!*

$$P(A, B) = P(A)P(B|A)$$

$$P(A, B, C) = P(A)P(B|A)P(C|A, B)$$

**The Chain Rule:**

$$P(X_1, X_2, \dots, X_n) = P(X_1)P(X_2|X_1)P(X_3|X_1, X_2) \dots P(X_n|X_1, \dots, X_{n-1})$$

# Language Modeling: How to Estimate?

V1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

V2: Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

*Observation: V1 and V2 are equivalent!*

$$P(A, B) = P(A)P(B|A)$$

$$P(A, B, C) = P(A)P(B|A)P(C|A, B)$$

**The Chain Rule:** 
$$P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | X_1, X_2, \dots, X_{i-1})$$

$$P(X_1, X_2, \dots, X_n) = P(X_1)P(X_2|X_1)P(X_3|X_1, X_2) \dots P(X_n|X_1, \dots, X_{n-1})$$

# Language Modeling: How to Estimate?

V1: Compute  $P(w_1, w_2, w_3, w_4, w_5) = P(W)$

V2: Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

*Observation: Solving V2 give us V1!*

$$P(A, B) = P(A)P(B|A)$$

$P(A, B,$

LM version 1

LM version 2

$$\underline{P(X_1, X_2, \dots, X_n)} = \underline{P(X_1, X_2, \dots, X_{n-1})} \underline{P(X_n | X_1, \dots, X_{n-1})}$$

**The Ch**

$$P(X_1, X_2, \dots, X_n) = P(X_1)P(X_2|X_1)P(X_3|X_1, X_2)\dots P(X_n|X_1, \dots, X_{n-1})$$

# Language Modeling: How to Estimate?

Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

**Problem:** even the Web isn't large enough to enable good estimates of most phrases.

**A solution:** Estimate from shorter sequences.

**Chain-Rule:**

$$P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | X_1, X_2, \dots, X_{i-1})$$



# Language Modeling: How to Estimate?

Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

**Problem:** even the Web isn't large enough to enable good estimates of most phrases.

**A solution:** Estimate from shorter sequences.

**Chain-Rule:**

$$P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | X_1, X_2, \dots, X_{i-1})$$

**Markov Assumption:**

$$P(X_n | X_1, \dots, X_{n-1}) \approx P(X_n | X_{n-k}, \dots, X_{n-1}) \quad \text{where } k < n$$

# Language Modeling: How to Estimate?

Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

**Problem:** even the Web isn't large enough to enable good estimates of most phrases.

**A solution:** Estimate from shorter sequences.

**Chain-Rule:**

$$P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | X_1, X_2, \dots, X_{i-1})$$

**Markov Assumption:**

$$P(X_n | X_1, \dots, X_{n-1}) \approx P(X_n | X_{n-k}, \dots, X_{n-1}) \quad \text{where } k < n$$

# Language Modeling: How to Estimate?

Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

**Unigram Model:  $k = 0$ ;** 
$$P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i)$$

# Language Modeling: How to Estimate?

Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

**Bigram Model:  $k = 1$ ;** 
$$P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | X_{i-1})$$

# Language Modeling: How to Estimate?

Compute  $P(w_5 | w_1, w_2, w_3, w_4) = P(w_n | w_1, w_2, \dots, w_{n-1})$

**Bigram Model:  $k = 1$ ;** 
$$P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | X_{i-1})$$

Example generated sentence:


*outside, new, car, parking, lot, of, the, agreement, reached*

$P(X_1 = \text{"outside"}, X_2 = \text{"new"}, X_3 = \text{"car"}, \dots)$

$\approx P(X_1 = \text{"outside"}) * P(X_2 = \text{"new"} | X_1 = \text{"outside"}) * P(X_3 = \text{"car"} | X_2 = \text{"new"}) * \dots$


# Language Modeling

Building a model (or system / API) that can answer the following:




# Language Modeling

Building a model (or system / API) that can answer the following:



# Language Modeling

Building a model (or system / API) that can answer the following:





# Language Model

## Building a model (

*a sequence of  
natural language*

Food corpus from Jurafsky (2018). Samples:

*can you tell me about any good cantonese restaurants close by*

*mid priced thai food is what i'm looking for*

*tell me about chez panisse*

*can you give me a listing of the kinds of food that are available*

*i'm looking for a good place to eat breakfast*

*when is caffe venezia open during the day*

Training Corpus

training  
(fit, learn)

## Bigram Counts

first word \ second word

	i	want	to	eat	chinese	food	lunch	spend
i	5	827	0	9	0	0	0	2
want	2	0	608	1	6	6	5	1
to	2	0	4	686	2	0	6	211
eat	0	0	2	0	16	2	42	0
chinese	1	0	0	0	0	82	1	0
food	15	0	15	0	1	4	0	0
lunch	2	0	0	0	0	1	0	0
spend	1	0	1	0	0	0	0	0

Example from (Jurafsky, 2017)

Training Corpus

training  
(fit, learn)

# Bigram Counts

first word | second word

	i	want	to	eat	chinese	food	lunch	spend
i	5	827	0	9	0	0	0	2
want	2	0	608	1	6	6	5	1
to	2	0	4	686	2	0	6	211
eat	0	0	2	0	16	2	42	0
chinese	1	0	0	0	0	82	1	0
food	15	0	15	0	1	4	0	0
lunch	2	0	0	0	0	1	0	0
spend	1	0	1	0	0	0	0	0

i	want	to	eat	chinese	food	lunch	spend
2533	927	2417	746	158	1093	341	278

Training Corpus

training  
(fit, learn)

# Bigram Counts

first word | second word

	i	want	to	eat	chinese	food	lunch	spend
i	5	827	0	9	0	0	0	2
want	2	0	608	1	6	6	5	1
to	2	0	4	686	2	0	6	211
eat	0	0	2	0	16	2	42	0
chinese	1	0	0	0	0	82	1	0
food	15	0	15	0	1	4	0	0
lunch	2	0	0	0	0	1	0	0
spend	1	0	1	0	0	0	0	0

i	want	to	eat	chinese	food	lunch	spend
2533	927	2417	746	158	1093	341	278

training corpus (fit, learn)

**Bigram model:**  $P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | X_{i-1})$

Need to estimate:  $P(X_i | X_{i-1}) = \text{count}(X_{i-1} X_i) / \text{count}(X_{i-1})$

second word:  $x_i$

$$P(X_i | X_{i-1})$$

first word:  $x_{i-1}$

	i	want	to	eat	chinese	food	lunch	spend
i	0.002	0.33	0	0.0036	0	0	0	0.00079
want	0.0022	0	0.66	0.0011	0.0065	0.0065	0.0054	0.0011
to	0.00083	0	0.0017	0.28	0.00083	0	0.0025	0.087
eat	0	0	0.0027	0	0.021	0.0027	0.056	0
chinese	0.0063	0	0	0	0	0.52	0.0063	0
food	0.014	0	0.014	0	0.00092	0.0037	0	0
lunch	0.0059	0	0	0	0	0.0029	0	0
spend	0.0036	0	0.0036	0	0	0	0	0

i	want	to	eat	chinese	food	lunch	spend
2533	927	2417	746	158	1093	341	278

Training corpus


**Bigram model:**  $P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | X_{i-1})$

Need to estimate:  $P(X_i | X_{i-1}) = \text{count}(X_{i-1} X_i) / \text{count}(X_{i-1})$

first word( $X_{i-1}$ ) \ second word ( $X_i$ )

$$P(X_i | X_{i-1})$$


	i	want	to	eat	chinese	food	lunch	spend
i	0.002	0.33	0	0.0036	0	0	0	0.00079
want	0.0022	0	0.66	0.0011	0.0065	0.0065	0.0054	0.0011



Need to estimate:  $P(X_i | X_{i-1}) = \text{count}(X_{i-1} X_i) / \text{count}(X_{i-1})$

i	0.002	0.33	0	0.0036	0	0	0	0.00079
want	0.0022	0	0.66	0.0011	0.0065	0.0065	0.0054	0.0011
to	0.00083	0	0.0017	0.28	0.00083	0	0.0025	0.087
eat	0	0	0.0027	0	0.021	0.0027	0.056	0
chinese	0.0063	0	0	0	0	0.52	0.0063	0
food	0.014	0	0.014	0	0.00092	0.0037	0	0

$$p(X_i = w \mid X_{i-1} = \text{food})$$




0
0
0
0
0
spend
278

count( $X_{i-1}$ )

# Language Modeling

Building a model (or system / API) that can answer the following:



Example from (Jurafsky, 2017)

training


⌈ **Bigram model:**  $P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | X_{i-1})$

Need to estimate:  $P(X_i | X_{i-1}) = \text{count}(X_{i-1} X_i) / \text{count}(X_{i-1})$



# Language Modeling

Building a model (or system / API) that can answer the following:



**Trigram model:** 
$$P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | X_{i-2}, X_{i-1})$$

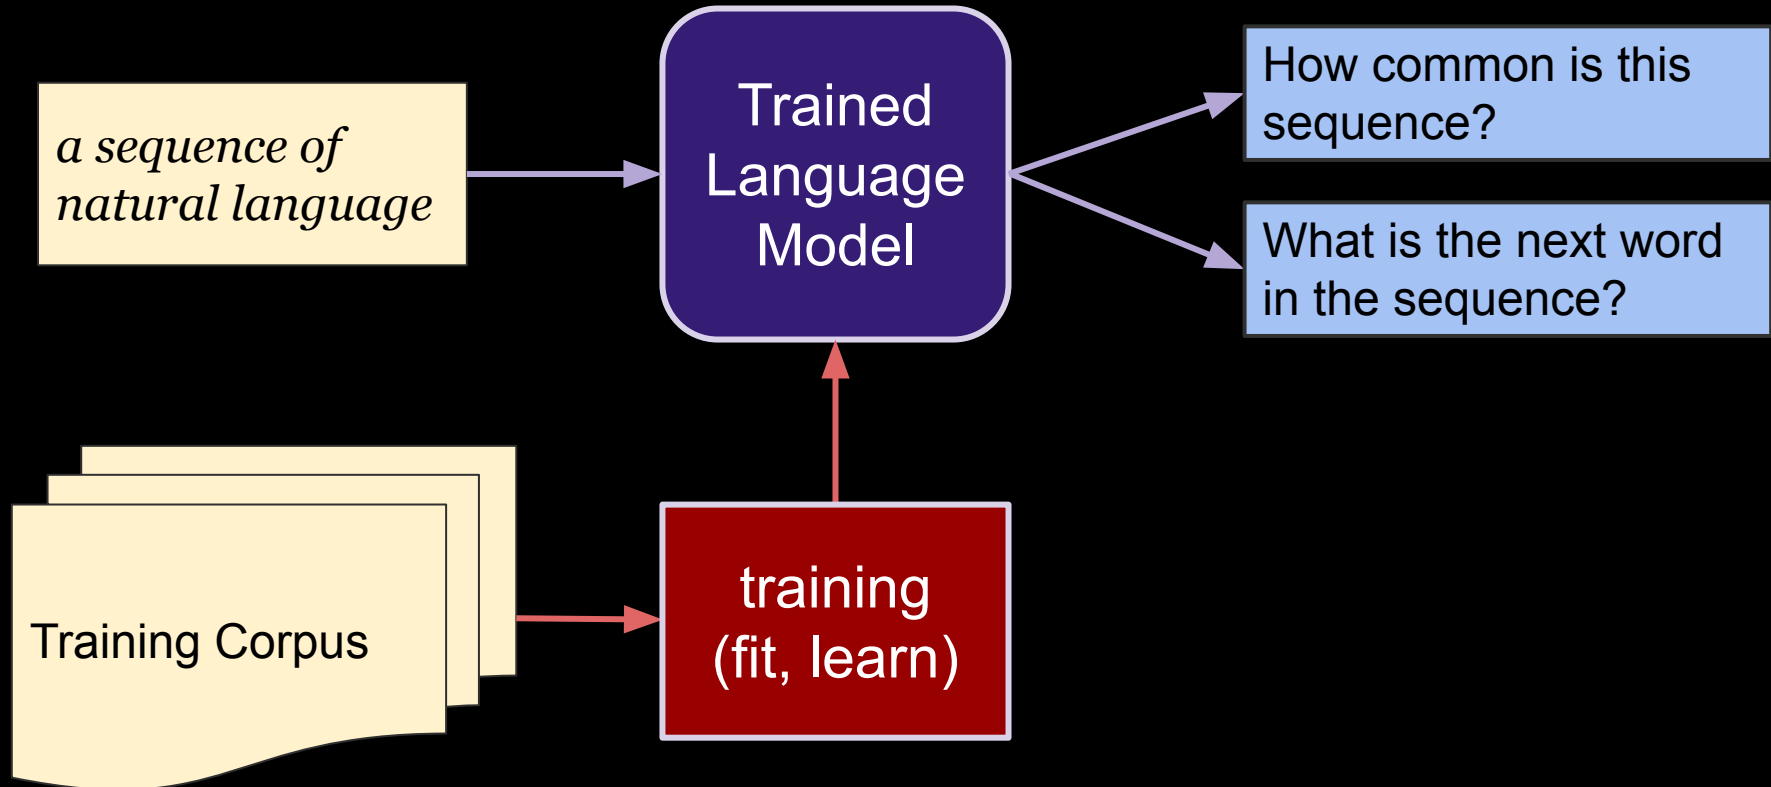
Need to estimate:  $P(X_i | X_{i-1}, X_{i-2}) = \text{count}(X_{i-2} X_{i-1} X_i) / \text{count}(X_{i-2} X_{i-1})$

**Bigram model:** 
$$P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | X_{i-1})$$

Need to estimate:  $P(X_i | X_{i-1}) = \text{count}(X_{i-1} X_i) / \text{count}(X_{i-1})$

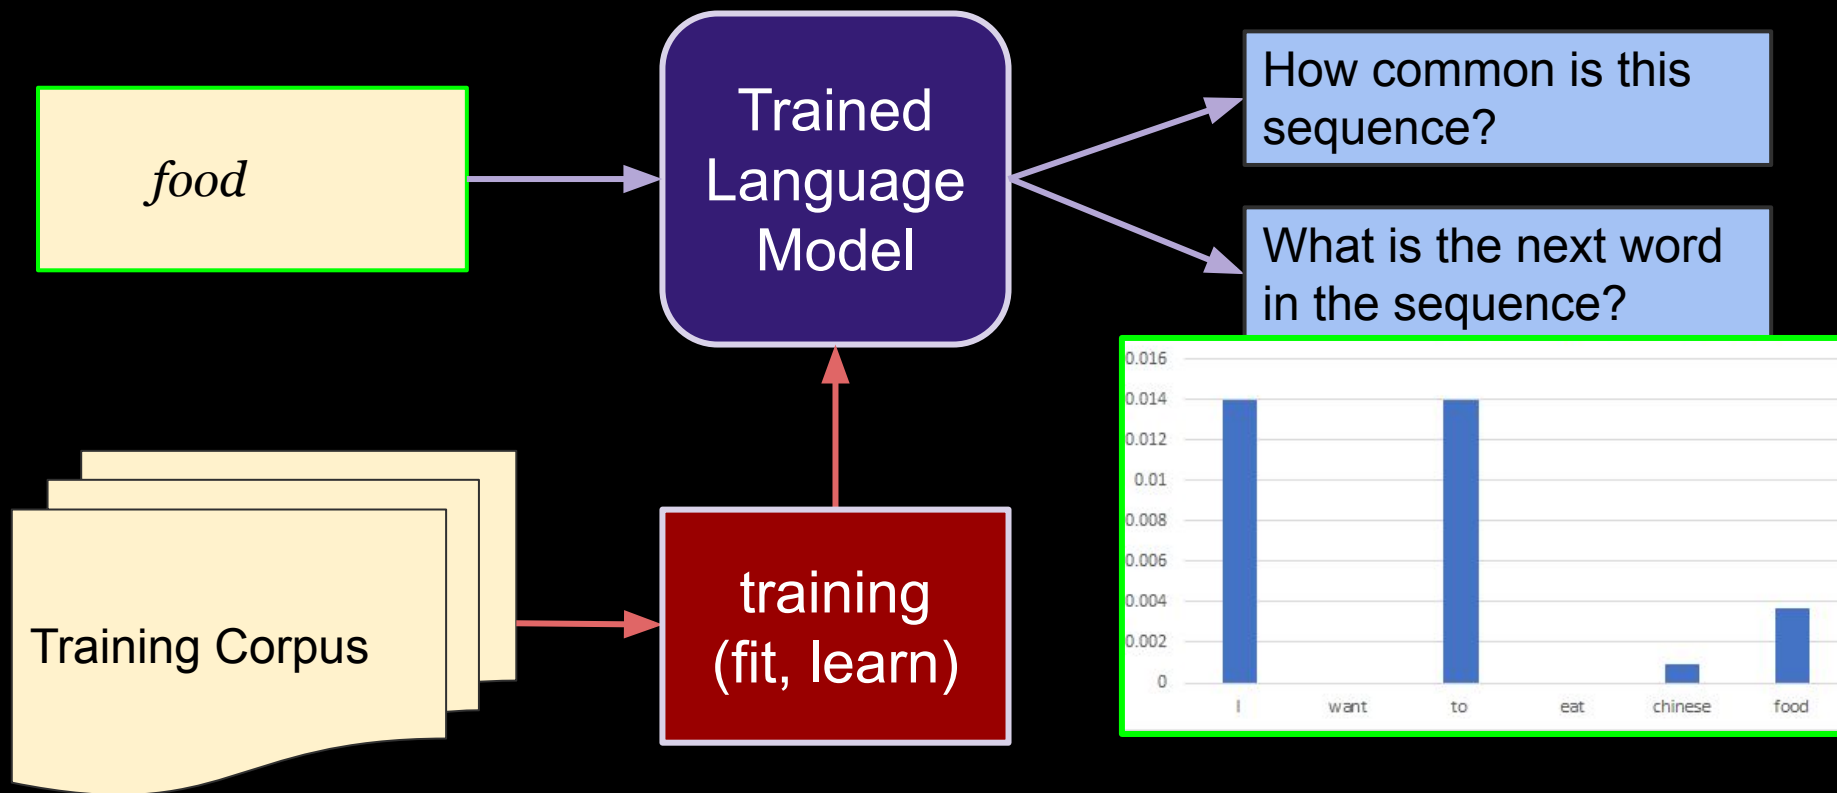
# Language Modeling

Building a model (or system / API) that can answer the following:




# Language Modeling

Building a model (or system / API) that can answer the following:




# Language Modeling

Building a model (or system / API) that can answer the following:




# Language Modeling

Building a model (or system / API) that can answer the following:




# Language Modeling


Building a model (or system / API) that can answer the following:



# Evaluation



# Evaluation




Apply Chain Rule:  $PP(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_1 \dots w_{i-1})}}$

$$PP(W) = P(w_1 w_2 \dots w_N)^{-\frac{1}{N}}$$
$$= \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}}$$



# Evaluation



Apply Chain Rule:  $PP(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_1 \dots w_{i-1})}}$


Thus,  
PP for Bigrams:

$$PP(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_{i-1})}}$$
$$PP(W) = P(w_1 w_2 \dots w_N)^{-\frac{1}{N}}$$
$$= \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}}$$

# Evaluation

Reasoning:

- 1) Inverse of probability  
(i.e. minimize perplexity = maximize likelihood)
- 2) (weighted) average branching factor



Apply Chain Rule:  $PP(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_1 \dots w_{i-1})}}$

Thus,  
PP for Bigrams:

$$PP(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_{i-1})}}$$
$$PP(W) = P(w_1 w_2 \dots w_N)^{-\frac{1}{N}}$$
$$= \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}}$$

# Evaluation

Reasoning:

- 1) Inverse of probability  
(i.e. minimize perplexity = maximize likelihood)
- 2) (weighted) average branching factor

**Qualitatively: Prefers real sentences**

(sequences that are more grammatical, make sense).

Test Corpus

Language  
Model

In the ... se?

Perplexity

Apply Chain Rule:  $PP(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_1 \dots w_{i-1})}}$

Thus,  
PP for Bigrams:

$$PP(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_{i-1})}}$$
$$PP(W) = P(w_1 w_2 \dots w_N)^{-\frac{1}{N}}$$
$$= \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}}$$

# Evaluation Summary

- Use *training set* to "learn model"  
(i.e. to store counts, from which we can derive probability for any  $p(w_i | w_{i-1}, w_{i-2})$ )
- Use held-out *testing set* to evaluate
- Perplexity -- metric for scoring how well learned model works on test.  
(an *intrinsic* evaluation)

Training 38 million words, test 1.5 million words, WSJ

N-gram Order	Unigram	Bigram	Trigram
Perplexity	962	170	109

# Evaluation Summary

- Use *training set* to "learn model"  
(i.e. to store counts, from which we can derive probability for any  $p(w_i | w_{i-1}, w_{i-2})$ )
- Use held-out *testing set* to evaluate
- Perplexity -- metric for scoring how well learned model works on test.  
(an *intrinsic* evaluation)
- *Extrinsic* evaluation: Test on task accuracies
  - machine translation: does it improve translation accuracy
  - autocomplete: do users like the suggestions
  - speech recognition: does it improve transcription accuracy
  - spelling corrector,  
etc...

## Practical Considerations for LMs:

- Use log probability for assessing perplexity to keep numbers reasonable and save computation.  
(uses addition rather than multiplication)
- Use Out-of-vocabulary (OOV)  
Choose minimum frequency or total vocabulary size and mark as <OOV>
- Sentence start and end: <*s*> *this is a sentence* </*s*>  
Advantage: models word probability at beginning or end.